


MASTERING LLM PRESENTS

COFFEE BREAK CONCEPTS



Caching Methods in Large Language Models (LLMs)




Follow us on



What is caching?


- 1 Store and reuse responses
- 2 Improve response time
- 3 Reduce cost



Types of Caching

- 1 In-memory caching
- 2 Disk-based caching
- 3 Semantic caching

Caching Strategies for LLM Optimization



In-Memory Caching


1

Stores data in RAM for rapid access.

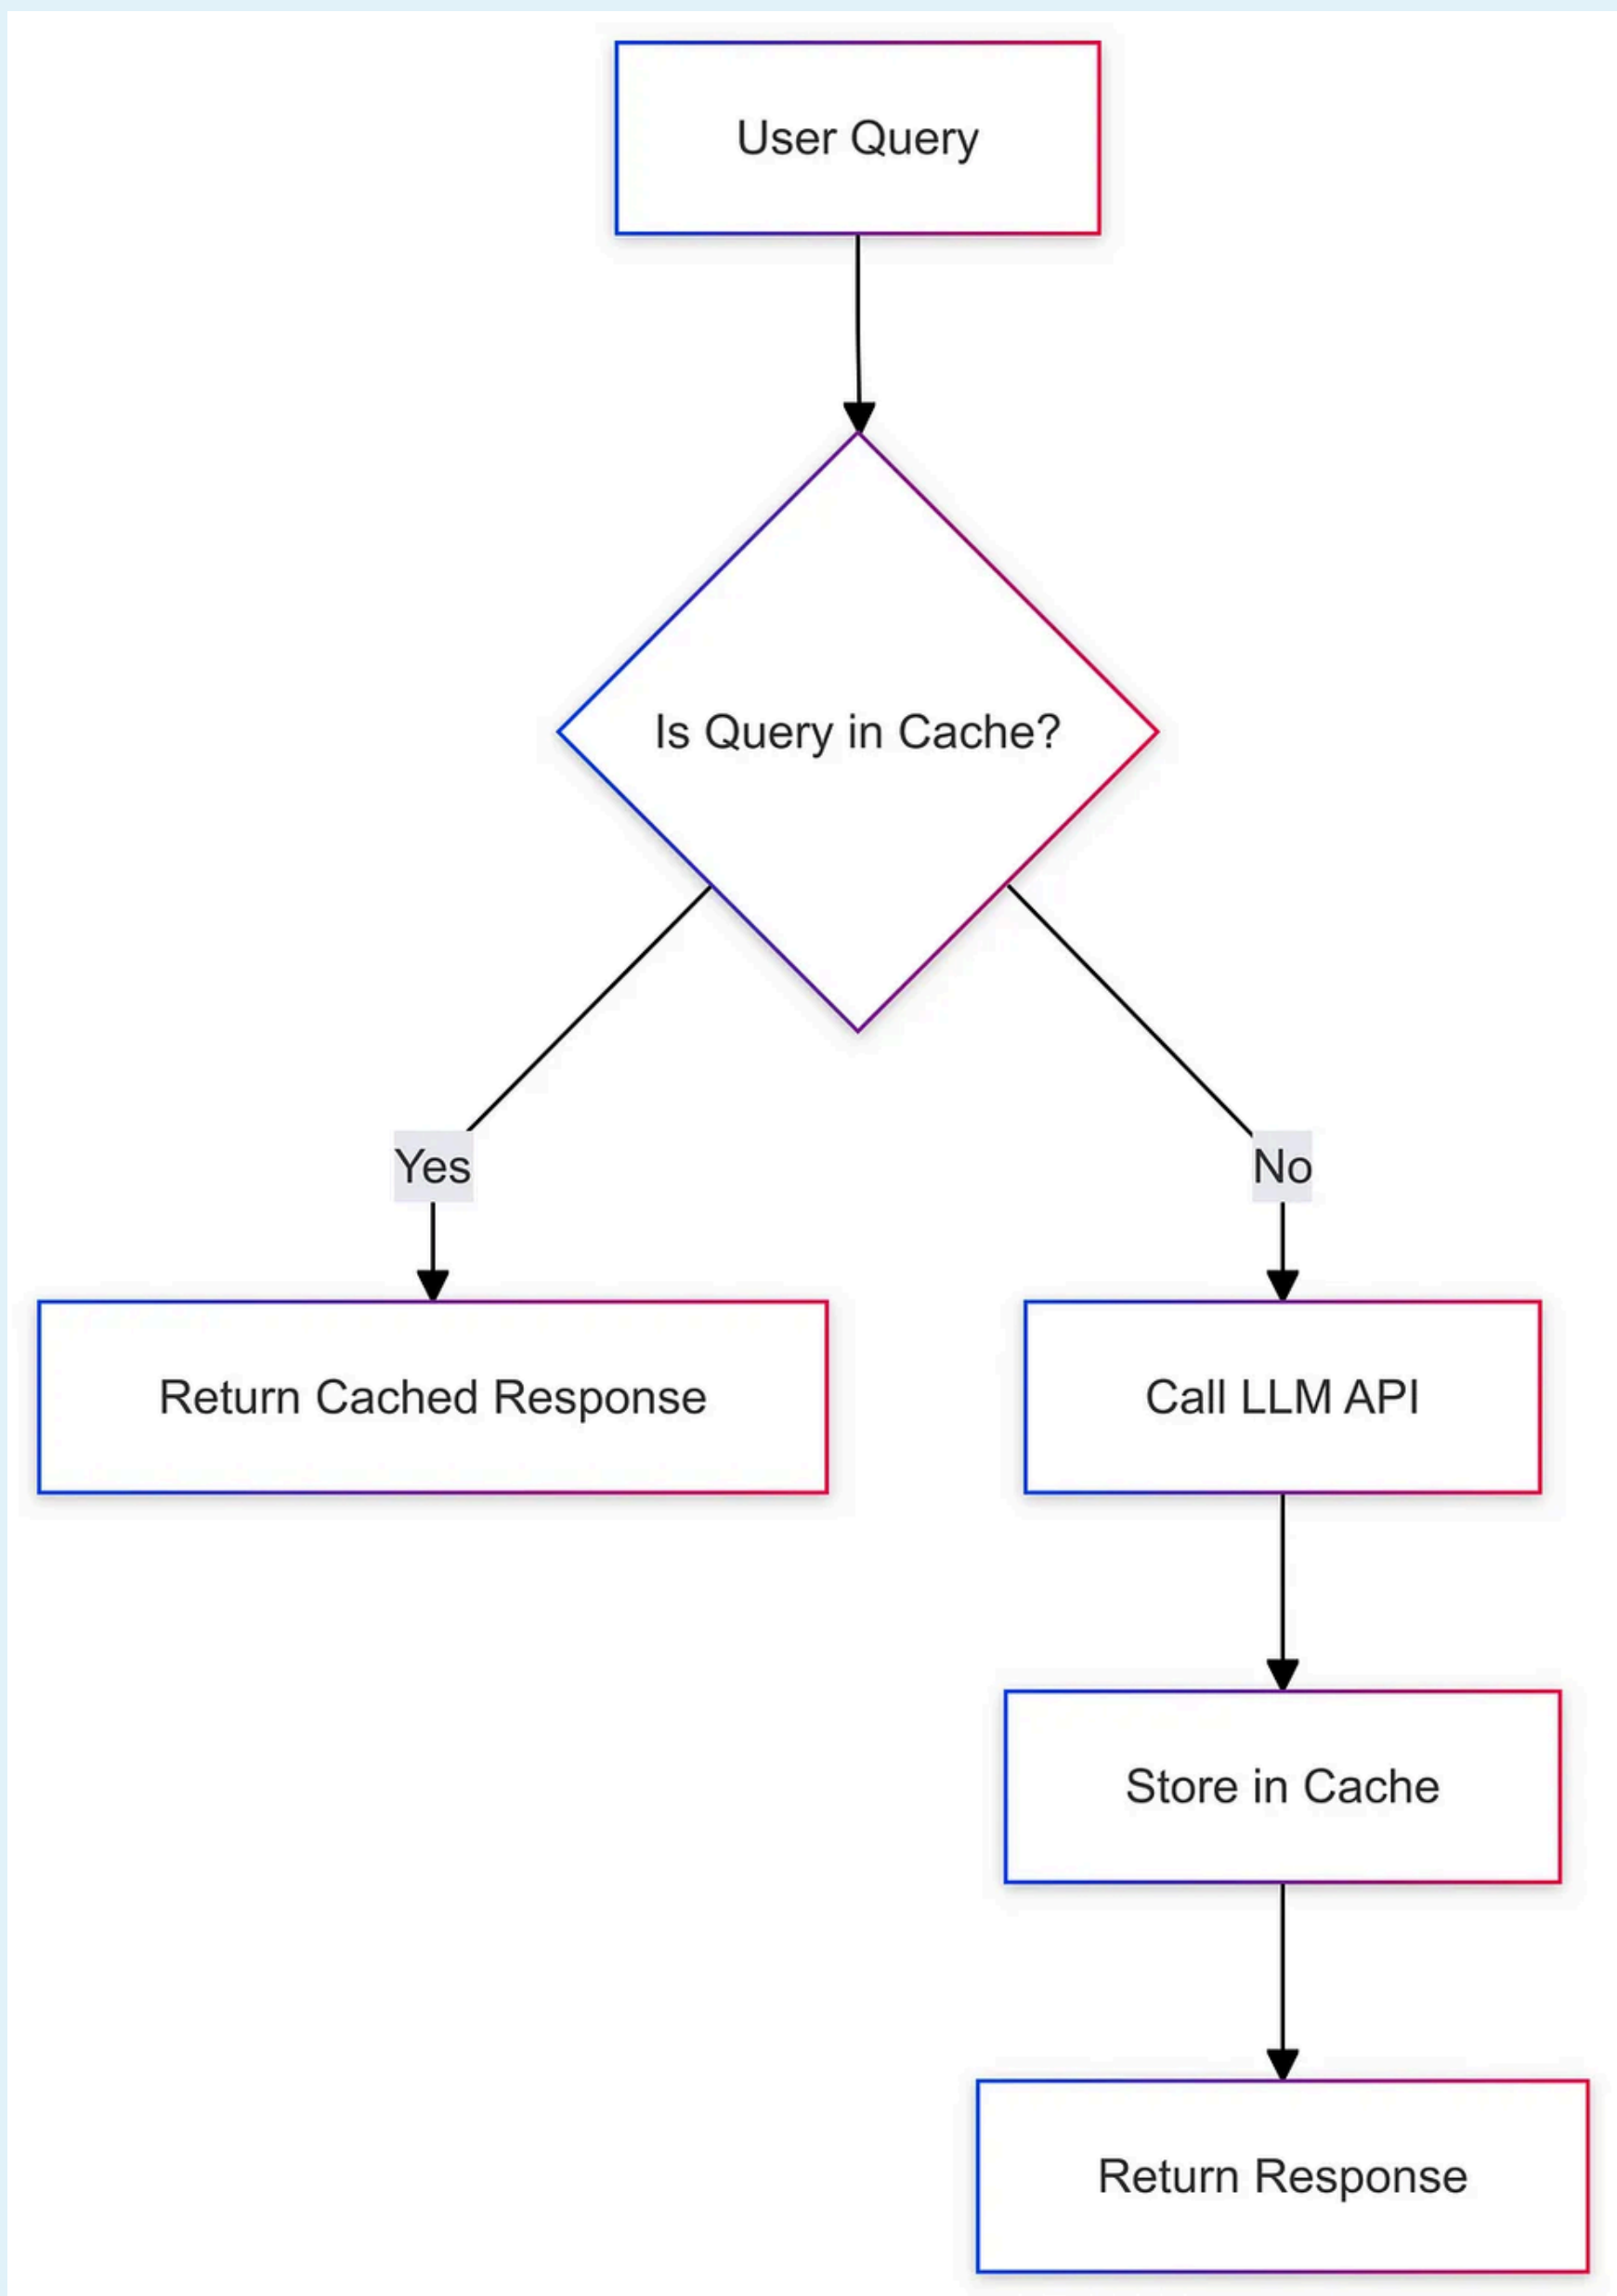
Example: Using a

2

dictionary in Python or tools like Redis for quick data retrieval.




In-Memory Caching



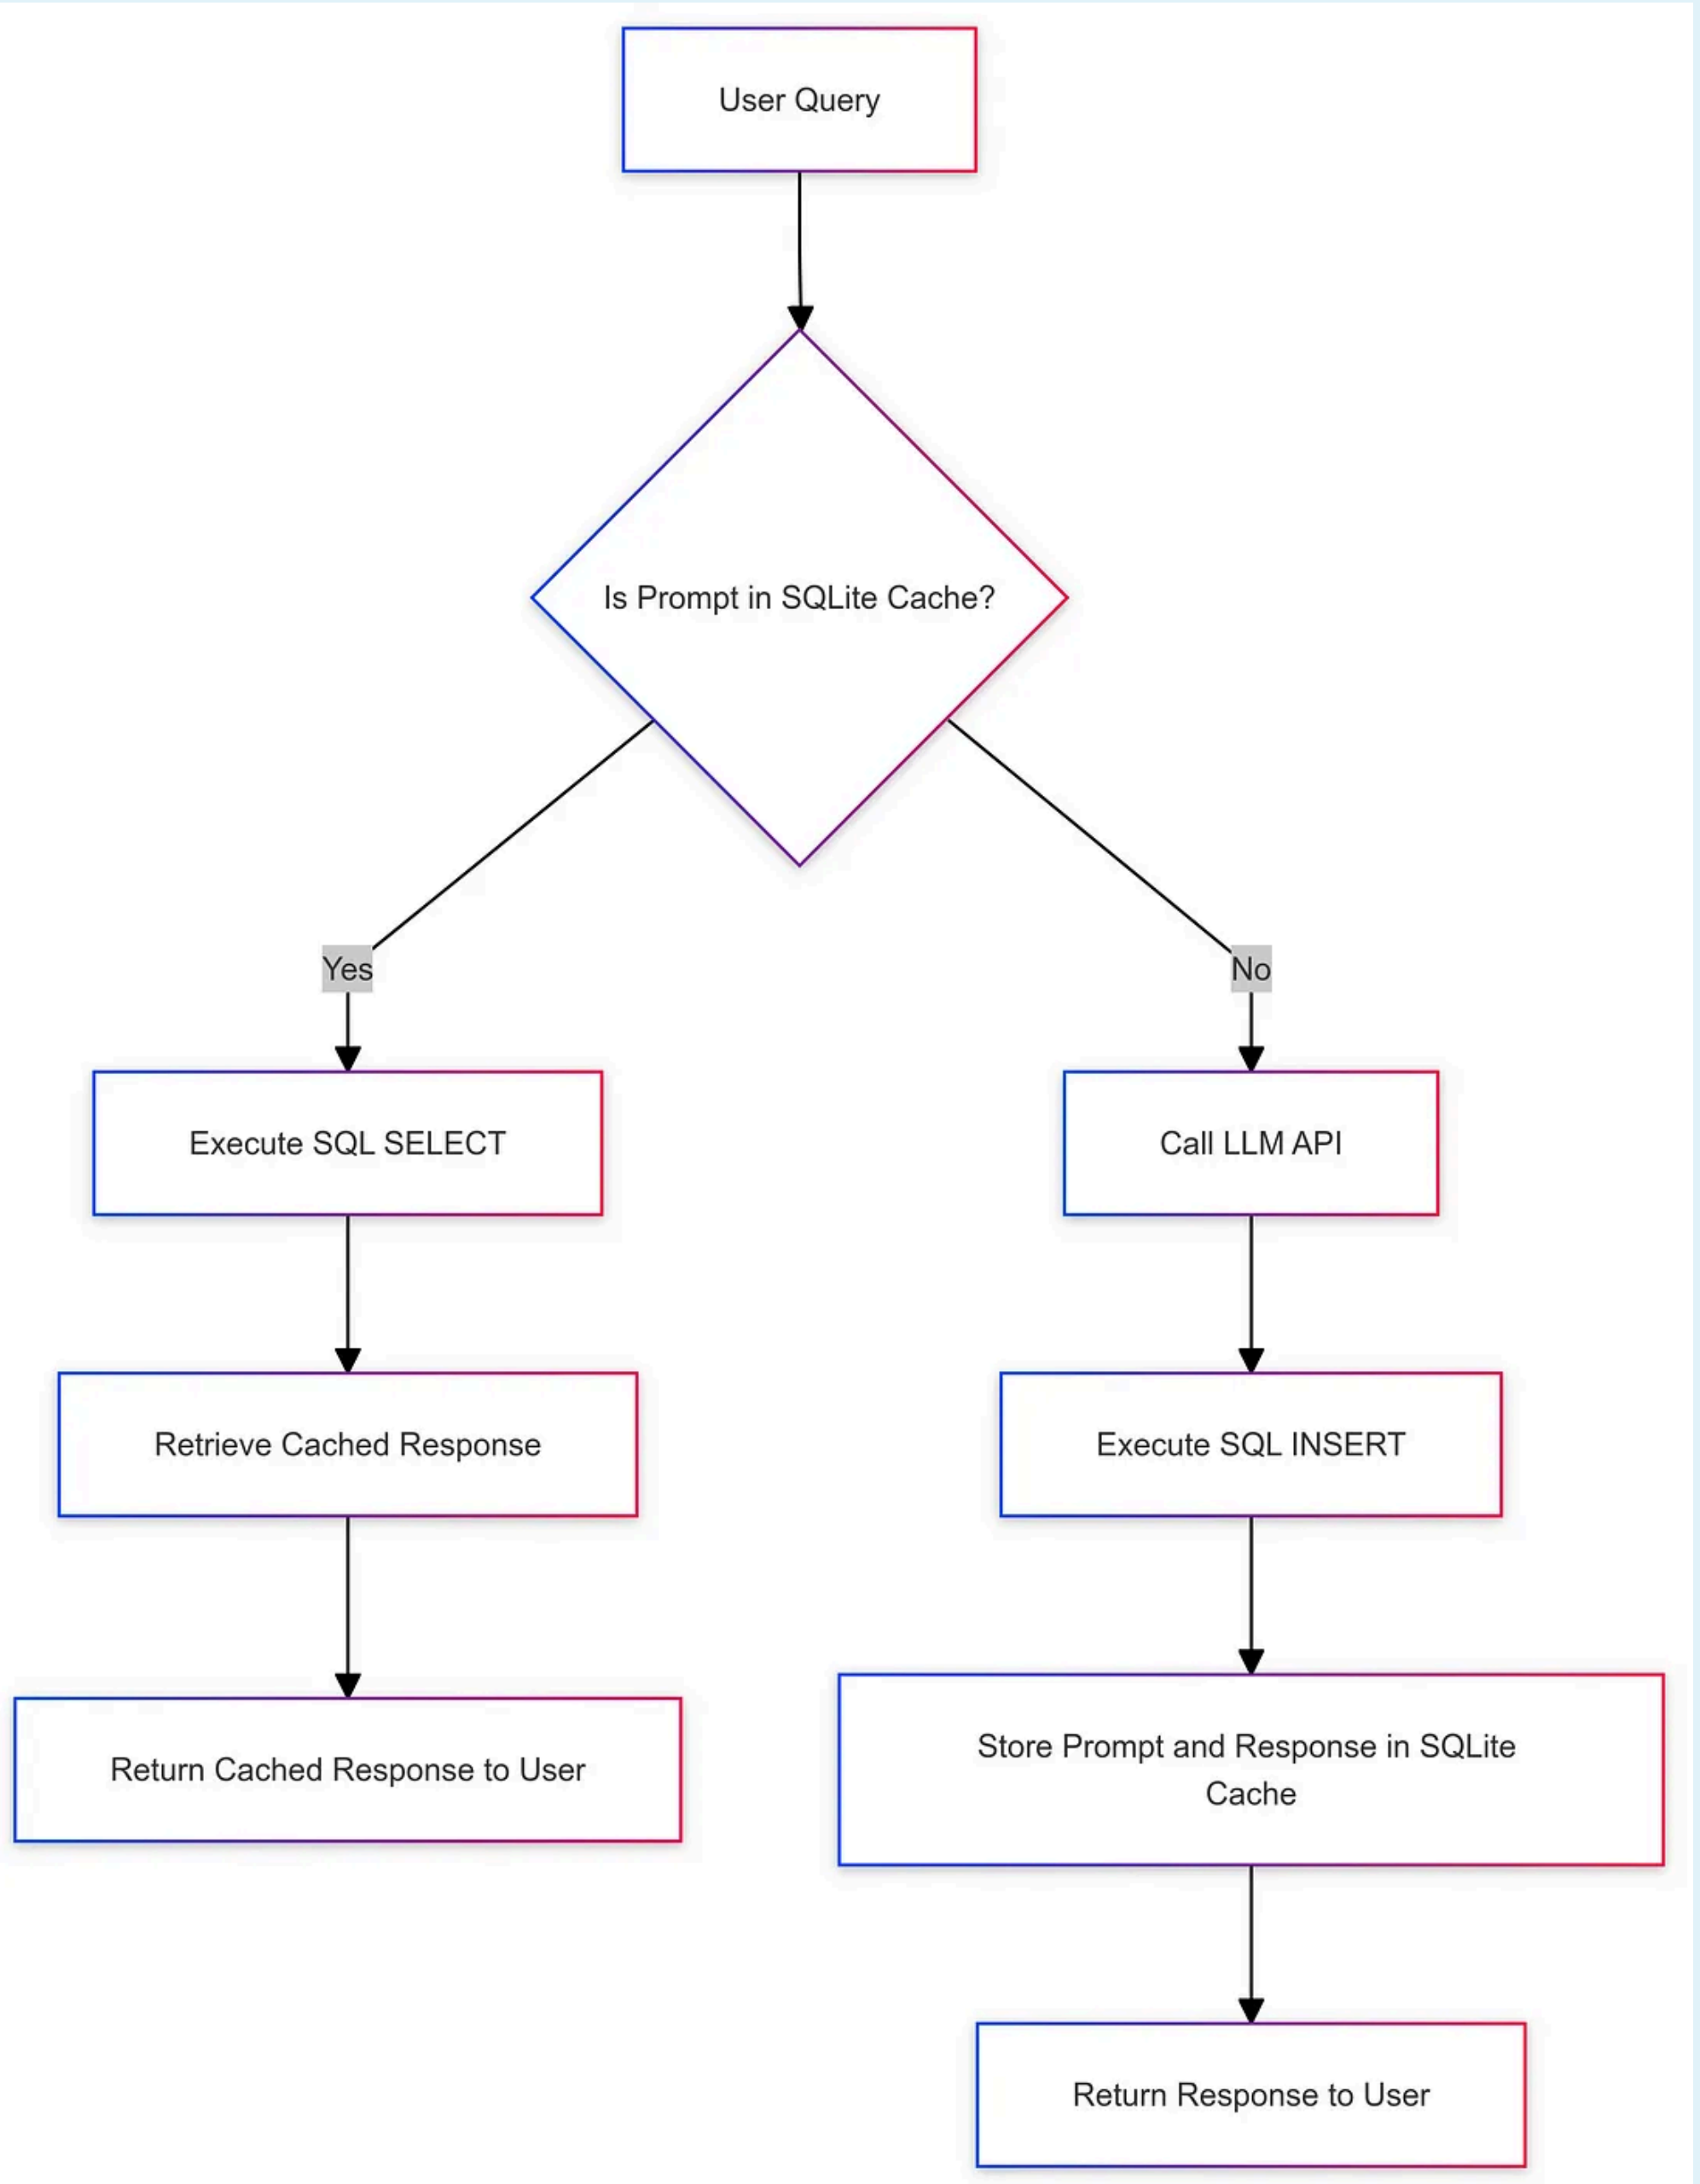
Disk-Based Caching

1 Stores data on disk using databases like SQLite.

2 **Example:** Utilizing SQLite or other disk-based databases to store cached responses.




Disk-Based Caching



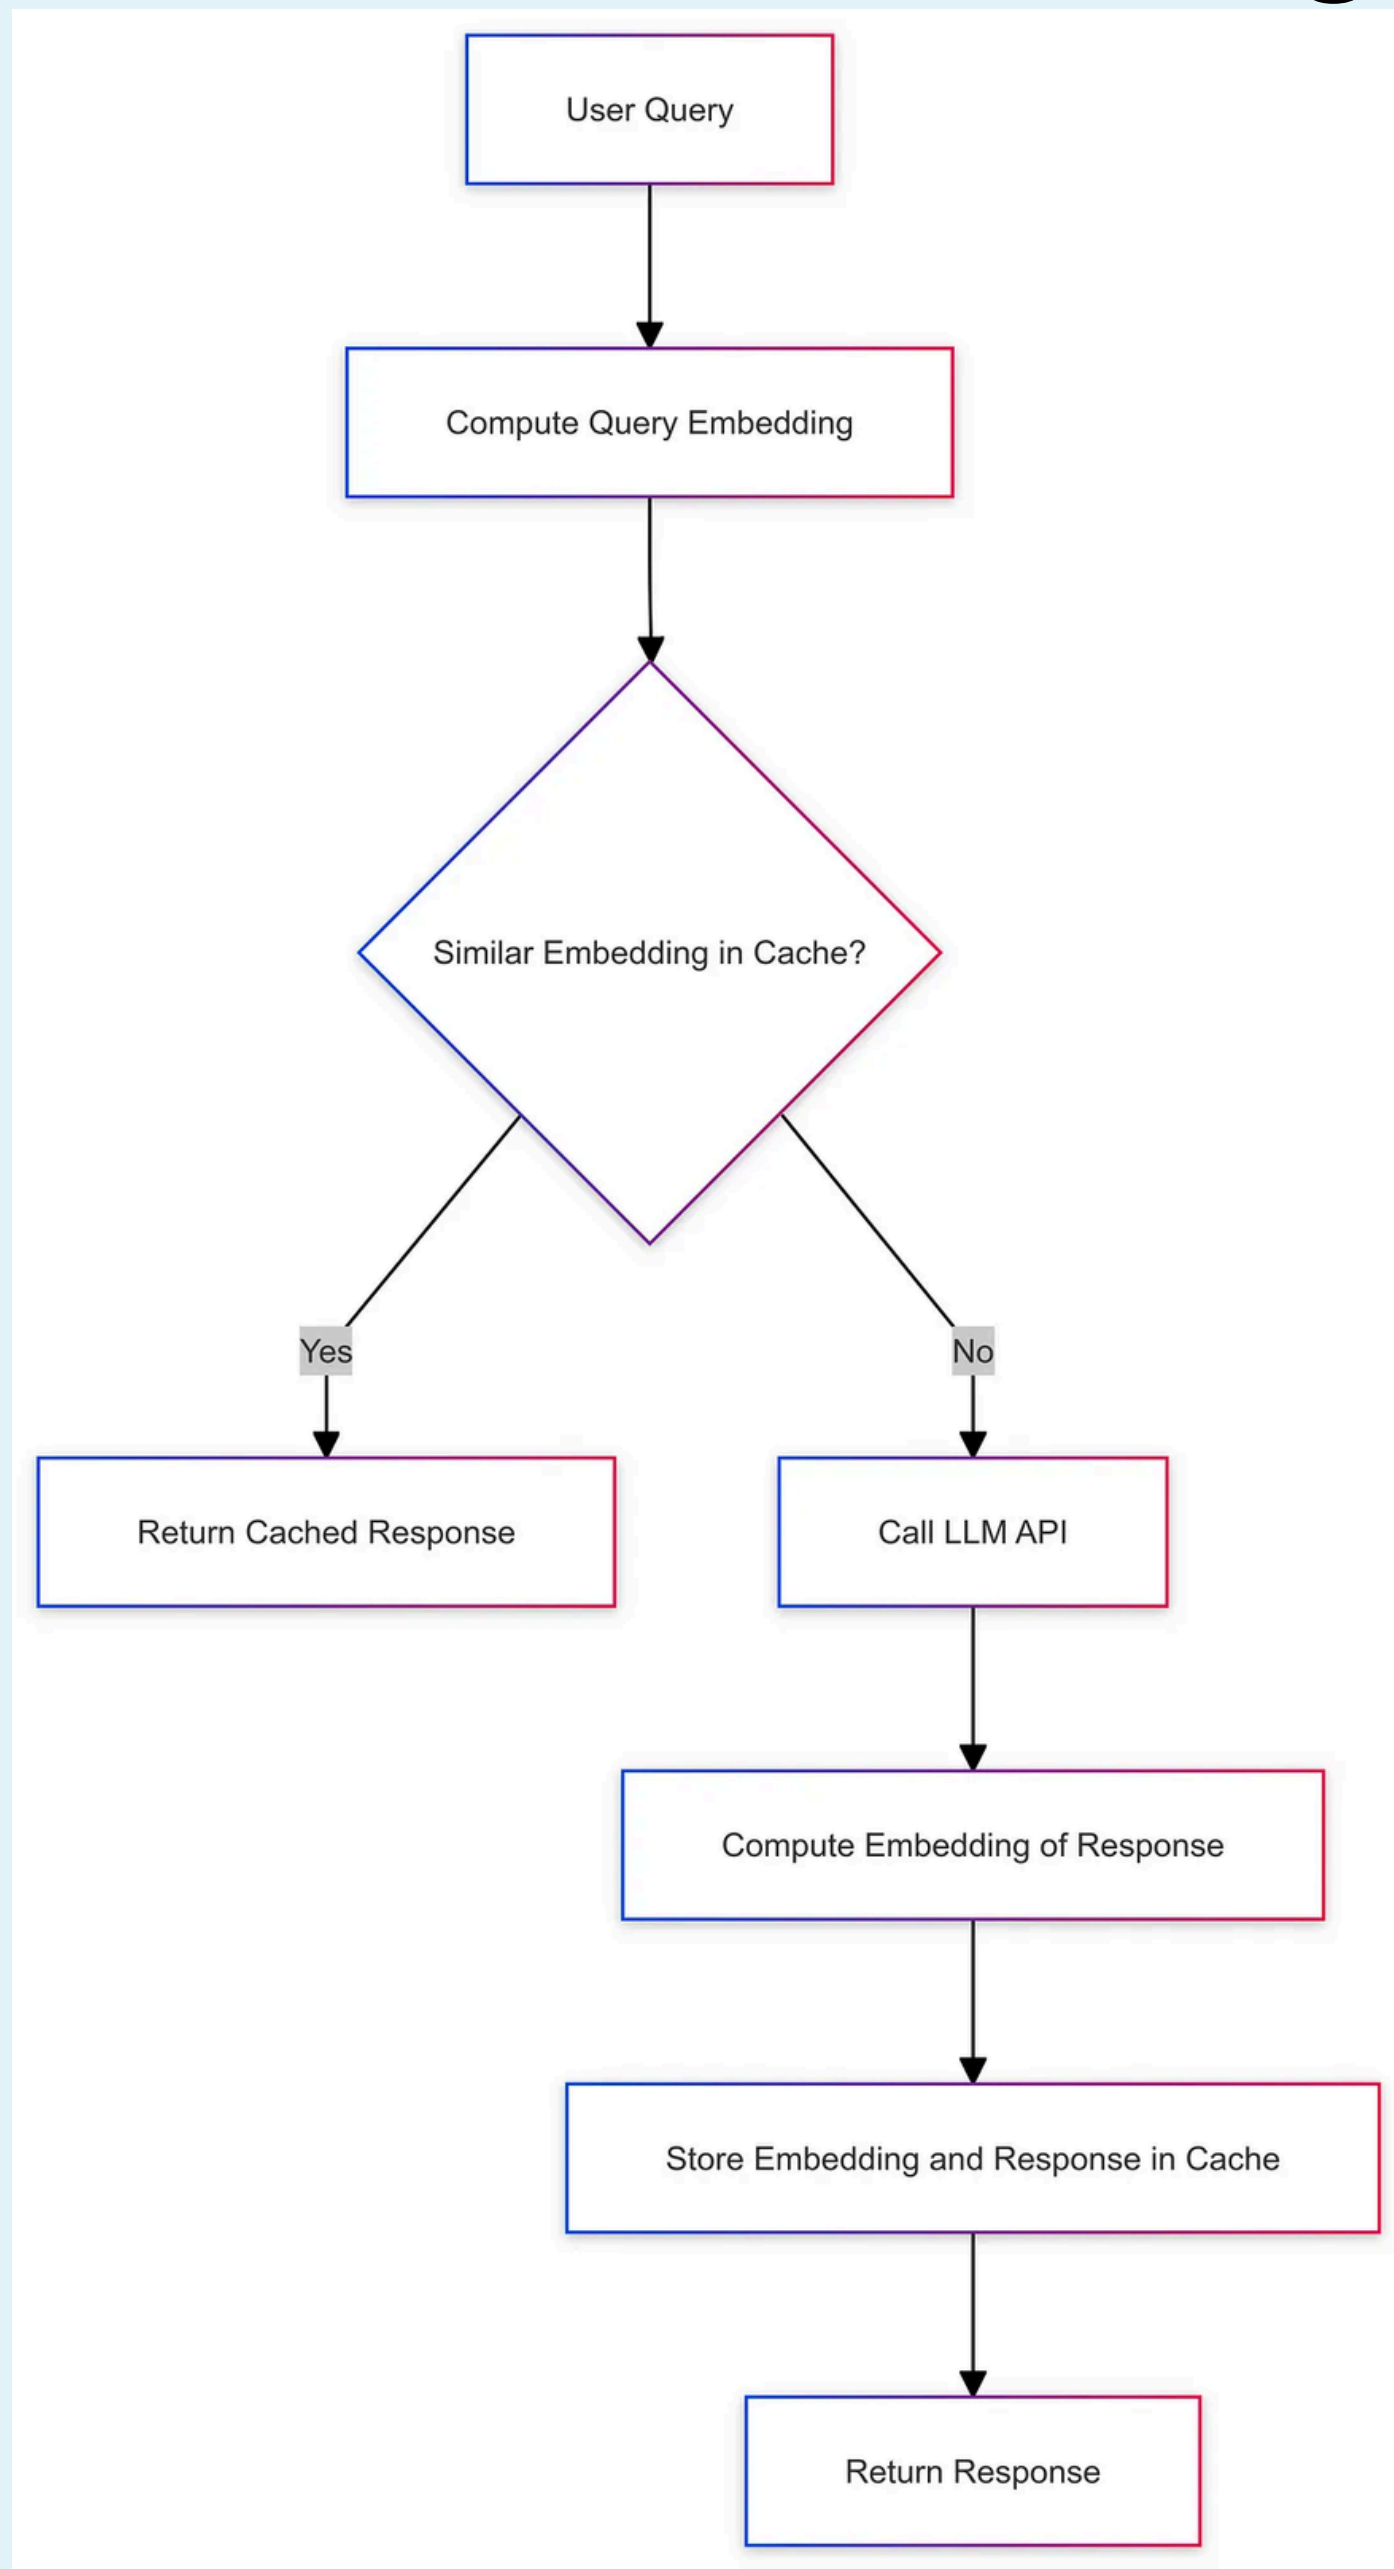
Semantic Caching

1

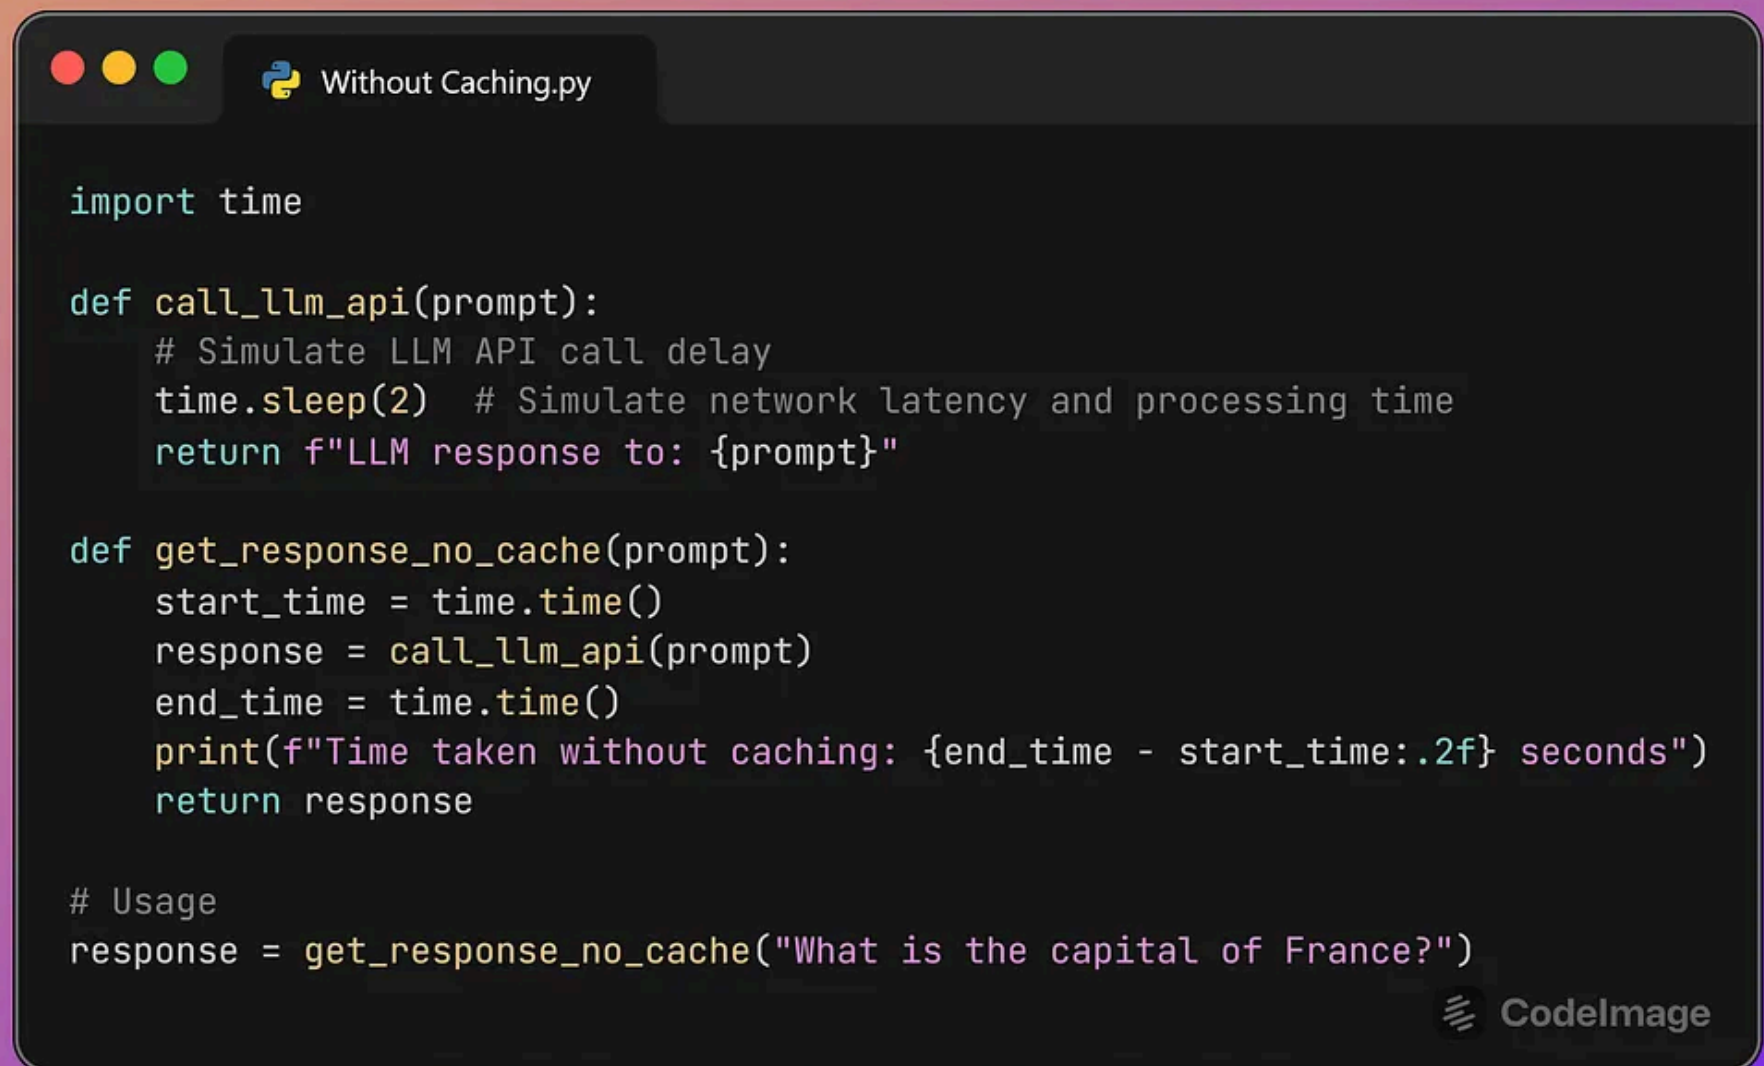
Stores responses based on the semantic meaning of queries using embeddings.



Semantic Caching



Without Caching



```
import time

def call_llm_api(prompt):
    # Simulate LLM API call delay
    time.sleep(2) # Simulate network latency and processing time
    return f"LLM response to: {prompt}"

def get_response_no_cache(prompt):
    start_time = time.time()
    response = call_llm_api(prompt)
    end_time = time.time()
    print(f"Time taken without caching: {end_time - start_time:.2f} seconds")
    return response

# Usage
response = get_response_no_cache("What is the capital of France?")
```

Output

```
Time taken without caching: 2.00 seconds
```

Note: We are not doing LLM call for simplicity.

Disk-Based Caching

```
import sqlite3
import time

def initialize_cache():
    conn = sqlite3.connect('cache.db')
    cursor = conn.cursor()
    cursor.execute('''
        CREATE TABLE IF NOT EXISTS cache (
            prompt TEXT PRIMARY KEY,
            response TEXT
        )
    ''')
    conn.commit()
    conn.close()

def call_llm_api(prompt):
    # Simulate LLM API call delay
    time.sleep(2) # Simulate network latency and processing time
    return f"LLM response to: {prompt}"

def get_response_disk_cache(prompt):
    start_time = time.time()
    conn = sqlite3.connect('cache.db')
    cursor = conn.cursor()
    # Check if the prompt exists in the cache
    cursor.execute('SELECT response FROM cache WHERE prompt=?', (prompt,))
    result = cursor.fetchone()
    if result:
        # Cache hit
        conn.close()
        end_time = time.time()
        print("Cache hit - Retrieved from SQLite")
        print(f"Time taken with disk-based caching (cache hit): {end_time - start_time:.2f} seconds")
        return result[0]
    else:
        # Cache miss
        response = call_llm_api(prompt)
        # Store the new prompt and response in the cache
        cursor.execute('INSERT INTO cache (prompt, response) VALUES (?, ?)', (prompt, response))
        conn.commit()
        conn.close()
        end_time = time.time()
        print("Cache miss - Not found in SQLite")
        print(f"Time taken with disk-based caching (cache miss): {end_time - start_time:.2f} seconds")
        return response

# Initialize cache database
initialize_cache()

# Usage
# First call (cache miss)
response = get_response_disk_cache("What is the tallest mountain in the world?")
print(response)

# Second call (cache hit)
response = get_response_disk_cache("What is the tallest mountain in the world?")
print(response)
```

Output

```
Cache miss - Not found in SQLite
Time taken with disk-based caching (cache miss): 2.05 seconds
LLM response to: What is the tallest mountain in the world?

Cache hit - Retrieved from SQLite
Time taken with disk-based caching (cache hit): 0.01 seconds
LLM response to: What is the tallest mountain in the world?
```

With Semantic Caching

```
import time
from sentence_transformers import SentenceTransformer
import faiss
import numpy as np

# Initialize the embedding model and index
model = SentenceTransformer('all-MiniLM-L6-v2')
embedding_dim = 384 # Embedding size for 'all-MiniLM-L6-v2'
index = faiss.IndexFlatL2(embedding_dim)
embeddings = []
responses = []

def call_llm_api(prompt):
    # Simulate LLM API call delay
    time.sleep(2) # Simulate network latency and processing time
    return f"LLM response to: {prompt}"

def get_response_semantic_cache(prompt, similarity_threshold=0.8):
    start_time = time.time()
    # Compute embedding for the prompt
    embedding = model.encode([prompt])[0]
    embedding = np.array([embedding]).astype('float32')

    if index.ntotal > 0:
        # Search for similar embeddings
        distances, indices = index.search(embedding, k=1)
        similarity = 1 - distances[0][0] / 4 # Normalize distance to similarity
        if similarity >= similarity_threshold:
            response = responses[indices[0][0]]
            print(f"Semantic cache hit (similarity: {similarity:.2f})")
        else:
            response = call_llm_api(prompt)
            index.add(embedding)
            responses.append(response)
            print("Semantic cache miss")
    else:
        # Cache is empty
        response = call_llm_api(prompt)
        index.add(embedding)
        responses.append(response)
        print("Semantic cache miss")
    end_time = time.time()
    print(f"Time taken with semantic caching: {end_time - start_time:.2f} seconds")
    return response

# First call (cache miss)
response = get_response_semantic_cache("Tell me about the Eiffel Tower.")

# Second call with a semantically similar prompt (cache hit)
response = get_response_semantic_cache("Give me information on the Eiffel Tower.")
```

CodeImage

Output

```
Semantic cache miss
Time taken with semantic caching: 2.35 seconds
Semantic cache hit (similarity: 0.89)
Time taken with semantic caching: 0.08 seconds
```


Tracking Performance

Cache Hit Ratio:

- 1 Percentage of requests served from the cache.
- 2 **Average Latency:** Time taken to serve requests.
- 3 **Cost Savings:** Reduction in API calls.

Summary

1

Choose Caching Type: Select In-Memory (fast, limited size), Disk-Based (persistent, larger capacity), or Semantic Caching (handles similar queries).

2

Optimize Cache Usage: Set clear cache rules; adjust similarity thresholds to balance hit rate and accuracy.

3

Leverage Tools: Use libraries like GPTCache; integrate with frameworks like LangChain or OpenAI API.

4

Monitor Performance: Track cache hits, latency, cost savings; refine strategy based on metrics.

www.masteringllm.com



LLM Interview Course



Want to Prepare yourself for an LLM Interview?

- ✓ 100+ Questions spanning 14 categories with Real Case Studies
- ✓ Curated 100+ assessments for each category
- ✓ Well-researched real-world interview questions based on FAANG & Fortune 500 companies
- ✓ Focus on Visual learning
- ✓ Certification



Coupon Code - LLM50

Coupon is valid till 30th Oct 2024

AgenticRAG with LlamaIndex

Want to learn why AgenticRAG is future of RAG?

- ✓ Master **RAG fundamentals** through practical case studies
- ✓ Understand how to overcome **limitations of RAG**
- ✓ Introduction to **AgenticRAG** & techniques like **Routing Agents, Query planning agents, Structure planning agents, and React agents with human in loop.**
- ✓ **5 real-time case studies with code walkthroughs**